


Day 1 — the gold-set pipeline, with real counts

Every number below is counted from a committed artifact, not typed in. The commands that reproduce each one are in the table at the bottom.

Read `docs/02_EVALUATION_SPEC.md` §1 for *why* the pipeline has these stages. This file is only concerned with *how many* rows survived each one, and *why the rest did not*.



The arithmetic that has to close

139 screened	
- 19 rejected on quality	(author's verdict: drop)
- 3 hand-excluded	(irreproducible, data/exclusions.json)
- 2 unvalidated	(stand-in screener, q115 and q117)
= 115 gold rows	(82.7% of screened)

Accepted rows are 110 `keep` + 5 `fix`. The 115 rows carry 150 chunk labels across 137 distinct chunks, because 26 rows were multi-labelled. The split lands at 80/35 — 69.6%, not exactly 70% — because whole groups of chunk-sharing rows have to move to the same side together.

Where the human sat in the loop

pass	date	rows	mode	agreement with the screener
disputed + controls	2026-08-16	62	revealed after each keypress	37/62
full coverage of agreed keeps	2026-08-17	73	fully blind	70/73 (96%)
rejection stratum (q017, q042)	2026-08-19	2	revealed	0/2
all decisions		137		107/137 (78%)

The 62-row and 2-row passes are the 64 "revealed" rows; 37/64 = 58%. Do not read 96% against 58% as an anchoring effect — see `DEFENCE.pdf` Q5.

Three notes the diagram cannot carry

- The drafting shortfall is not in an artifact.** 350 chunks were sampled and 148 draft rows exist. The spec explains the mechanism (Groq's daily token budget, `02_EVALUATION_SPEC.md` §1 "Budget"), but no run log, error row, or counter in this repository records where the run stopped. The "-202" in the diagram is `350 - 148`, not a measured figure.
- The eligible pool moved after drafting.** The scaffolding filter was added mid-Day-1 in response to the first audit. Nine drafted rows came from chunks that the filter now rejects. Three were excluded by hand for exactly that reason; five never made the gold set anyway; **one, q017, is in the gold set** and meets the stated exclusion criterion. See `DEFENCE.pdf` Q8.
- n 350 cannot be re-run to the same sample.** The eligible pool was 405 chunks when the drafting run sampled it and is 392 now, so `sample_chunks(eligible, 350, 42)` today draws a different 350. The gold set is reproducible from `goldset_screened.jsonl` forward; it is not reproducible from the corpus.

Reproducing each number

number	command
127 documents, 1,480 chunks	<code>python -m backend.ingest --corpus data/demo_corpus --dry-run</code>

number	command
392 eligible, and the four rejection counts	snippet below
148 drafted, 139 offered	<code>wc -l data/goldset_draft.jsonl data/goldset_draft_supplement.jsonl</code> , then count <code>status</code>
139 screened, keep/fix/drop 128/2/9	count <code>screen_verdict</code> in <code>data/goldset_screened.jsonl</code>
second opinion 90/49, disputed 50	join <code>data/claude_screen.jsonl</code> on <code>qid</code>
137 decisions, 73 blind, 107 agreed	<code>data/audit_results.json</code> , written by <code>backend.goldset.audit</code>
115 gold, 19 rejected, 82.7% accepted	<code>python -m backend.goldset.assemble</code> (prints the census)
26 multi-labelled rows, 35 extra chunks	<code>data/multilabel.json</code>
train 80 / test 35, no leakage	<code>python -m backend.goldset.split</code>

The eligible-pool counts come from `select_eligible`, which needs no database and no API key:

```

from pathlib import Path
from backend.ingest import build_chunks, get_tokenizer
from backend.goldset.generate import select_eligible

chunks = build_chunks(
    corpus_dir=Path("data/demo_corpus"), corpus_id="demo",
    tokenizer=get_tokenizer("BAAI/bge-small-en-v1.5"), size=512, overlap=64,
)
pool = select_eligible(chunks)
print(len(pool.eligible), pool.too_much_code, pool.code_edges,
      pool.scaffolding, pool.too_short)
# 392 748 325 14 1

```